

Responses to open questions

Draft for review (v3) · 28 September 2026 · Author Ben Fehnert

These responses draw on Kind’s working documents, investor materials, example reports and website. Each bullet shows its source in brackets (see the key below). **Notes** record assumptions, limits and open points. **Science Advisory Board input requested** marks areas where Kind is asking for the Board’s support. The final section lists inconsistencies found between the sources.

SUMMARY

Q	Topic	Status
1	Goals, purpose and claims	Partly answered
2	Scope of n-of-1 questions	Answered
3	Who decides safety and appropriateness	Answered
4	Design decisions and methodological framework	Answered
5	Outcome selection and interpretation	Science Advisory Board input requested
6	What the user is told	Science Advisory Board input requested
7	Safety monitoring and stopping rules	Partly answered
8	Regulatory framework	Answered
9	Data use, ownership and consent	Partly answered
10	Aggregating results across individuals	Partly answered

SOURCE KEY

BRIEF	Kind IRB brief, 21 Aug 2026 (author Ben Fehnert, reviewed by Prof Glyn Elwyn). § numbers refer to its assumptions (§4.x) and analysis types (§10a–e). It still names the Science Advisory Board as the approving body; this has since moved to the Protocol Review Board.
BRANY	Interim answers to the IRB’s open questions, 7 Sep 2026 (questions A–F). An analysis at this stage, to be built out and shared with legal counsel.
CAT	Kind Health Explorations Catalog, 28 Sep 2026 (v5): “Explorations Catalog update”, “Kind liability considerations” and “AI Matching Logic” tabs.
DECK	Kind Organization deck, 28 Sep 2026 (v5). “s” = slide number.
R&R	Kind Roles and Responsibilities, 28 Sep 2026 (v5). “p” = page number.
PITCH	Kind pitch deck (angel round), 21 Sep 2026. “p” = page number.
VB	Kind venture brief, 21 Sep 2026. “p” = page number.
REPORT	Example reports for the Morning rules exploration (“Anna”), 29 Jun 2026. “p” = page number.
WEB	Kind website, https://kind-health.app : home page, FAQ and privacy policy.

Q1 Goals, purpose and claims **Partly answered**

What are your goals for Kind? Is the product primarily a tool for personal decision-making, clinical care, research, wellness, or some combination—and what claims do you intend to make about its benefits?

- Kind helps an individual understand the extent to which a specific intervention relates to a change in outcome for them (intended use for launch). [DECK s4]
- It combines personal decision-making and wellness (individuals choose, track and see their own results) with research (opt-in, de-identified data for group analyses and preprints). It is not a clinical care tool: “not a medical service and does not provide medical advice”. [BRIEF §1–3; DECK s3; WEB FAQ]
- Business goals: a soft and full US launch in 2027 for the Optimizer and Improver segments, growing to about 1.2M US users by 2029, with research publishing from launch. Later phases add clinical explorations (2028) and global scale (2030). [VB p1–2; PITCH p14, p17]
- The launch scope is general wellness only (Tier 1, clinician-exempt), deliberately designed to sit outside FDA device oversight. [CAT liability tab, FDA; DECK s11]
- Product statement: “Not intended for use in the diagnosis, treatment, cure, mitigation, or prevention of any disease or medical condition; not intended for use by individuals to manage a diagnosed condition or make treatment decisions.” [CAT liability tab, Tier 1 product liability; DECK s4]

NOTE: The benefit claims Kind intends to make (in marketing or in the product) are not yet fully developed and will need to be worked out as part of product launch, making sure any claims are within the intended use statement. Some current website wording will need to be checked against it (see inconsistencies 1–3).

Q2 Scope of n-of-1 questions **Answered**

What kinds of n-of-1 questions will Kind accept, and what is explicitly out of bounds? In particular, will users experiment with prescription drugs, dose changes, drug discontinuation, supplements, diets, devices, or other interventions that could carry appreciable risk?

- Individuals choose from a curated catalog of pre-specified explorations; they do not design their own experiments. Adults 18+ only, with each exploration’s eligibility and exclusions applied as hard filters. [BRIEF §2, §4.5, §4.7; CAT update tab]
- At launch (Tier 1, 40 of 58 explorations): behaviour, sleep, diet, exercise, over-the-counter supplements (e.g. creatine, magnesium glycinate, L-theanine) and simple tracking. Most explorations run 2–8 weeks; no special equipment is needed, and wearable data is optional. [CAT update tab; DECK s15; WEB FAQ]
- Out of bounds at launch: prescription drugs, dose changes and discontinuation. Tier 3 medication comparisons come only after FDA review; they are prescriber-directed, mostly use sequential titration without withdrawing treatment, and Kind never suggests a dose or treatment change. [CAT header note, liability Tier 3; DECK s12, s16]
- Higher-risk areas are controlled: calorie restriction is Tier 2 with a safe intake floor; cannabinoids are observational only (an interventional design would need an FDA IND, plus DEA registration for THC); psilocybin has been removed as it is a Schedule I substance in the US. Active screening for vulnerable users (e.g. SCOFF before diet explorations) is proposed. [CAT update tab review flags; BRANY D]

NOTE: The boundaries are currently implied by the tier rules and individual catalog decisions, but Kind could develop a written policy on what is in and out of bounds.

Q3 Who decides safety and appropriateness **Answered**

Who determines whether a proposed experiment is clinically safe and scientifically appropriate? What clinical expertise is involved, when is the user’s own clinician brought in, and who has authority to reject a proposed experiment?

- Every exploration is reviewed by Kind’s Protocol Review Board and then approved by BRANY, an independent AAHRPP-accredited IRB, before launch. Researchers can review protocols they develop with Kind and follow the same review and IRB process. [DECK s9–10; R&R pp.8, 16–17; PITCH p11]

- The Protocol Review Board is chaired by the CSO, Prof Glyn Elwyn (MD), and run by the Head of Science. Its quorum includes a non-scientist and an unaffiliated member; Science Advisory Board members (e.g. Prof Sunita Vohra, MD) and health category specialists are invited as needed. It carries out a clinical and user-safety risk assessment, and the DPO signs off. [DECK s9; R&R pp.4, 8]
- Authority to reject: the Protocol Review Board can decline or send back a protocol; the DPO can withhold sign-off; BRANY decides whether an exploration fits an approved template and can require fuller review. [R&R p.11; BRIEF §13a; DECK s10]
- The user's clinician: not involved for Tier 1, though users with a diagnosed condition or medication are advised to consult a healthcare professional. Tier 2 requires the user to confirm they understand they should consult their clinician. Tier 3 requires an individual to confirm they have both a valid prescription and a prescriber relationship. [CAT liability tab, tier parameters; WEB FAQ]

NOTE: A practising clinician is not a fixed member of the Protocol Review Board; clinical input relies on the CSO and invited specialists. This will need to be confirmed and reviewed with BRANY, with input from the Science Advisory Board.

Q4 Design decisions and methodological framework Answered

How are the major design decisions made? How will you determine comparators, randomization, blinding, treatment periods, washout, dose, duration, adherence assessment, and outcome measures? Is there a standardized methodological framework?

- Yes: each exploration uses one of seven single-case archetypes chosen by reversibility, onset speed and safety of withdrawal (alternating treatments, ABAB reversal, changing criterion, sequential titration, split application, ABA elimination-reintroduction, observational). [CAT header note; DECK s16]
- Each catalog entry pre-specifies arms, duration, baseline period, outcomes and eligibility. Alternating designs randomise the order of days, with a minimum number of days per condition. Washout is handled with extended phases where effects persist (e.g. a 4-week withdrawal for creatine), and doses are fixed per exploration (e.g. 5 g/day creatine). [CAT update tab]
- Protocols are developed against SPENT/SPIRIT and CENT/CONSORT guidance and revised by the Protocol Review Board against approved templates, before being approved by the IRB prior to being added to Kind. [DECK s3, s9, s14; PITCH p8]
- Adherence is self-reported through daily logging, supported by AI companion reminders (typically once or twice a day). Reports show logging adherence (e.g. 92.9% of days) and which elements were followed each day. [WEB FAQ; VB p1; REPORT pp.1, 3]

NOTE: While blinding is possible, it is hard to manage and is currently only considered for one Tier 3 exploration (ADHD medication). For launch, we are assuming all adherence is self-reported. The current catalog notes that durations, outcomes and eligibility are still to be validated by the Protocol Review Board.

Q5 Outcome selection and interpretation Science Advisory Board input requested

How are outcomes selected and interpreted? Are validated measures preferred? How do you deal with measurement error, clinically meaningful change, multiple outcomes, missing observations, regression to the mean, carryover, and natural fluctuations in symptoms?

- Outcomes are pre-specified per exploration, mostly short self-reported ratings (e.g. afternoon energy 0–10, crash severity 0–3), with optional wearable data. [CAT update tab; REPORT p1]
- Natural fluctuation and carryover are partly addressed through design: baseline periods, randomised alternation, replicated blocks (ABAB) and washout phases. [CAT header note and design rationale]
- Reports already state the main limitations to users: outcomes are self-reported and open to expectation effects, other life changes may play a part, and without a return-to-baseline phase some change may reflect natural drift. [REPORT pp.4, 8]
- Group analyses only include individuals who meet data quality and completeness thresholds. [BRIEF §10b]

NOTE: While Kind does have a preference for validated patient-reported outcome measures (PROMs), these will be reviewed case by case, taking licensing, cost and ease of use into consideration. For launch, Kind will not use outcome measures that are disease indicators (e.g. PHQ-9, GAD-7).

SCIENCE ADVISORY BOARD INPUT REQUESTED: This is a key area where we are looking for support from the Science Advisory Board: to create a statistical analysis framework that takes account of measurement error, minimal clinically important change, multiple outcomes, missing data and regression to the mean.

Q6 What the user is told **Partly answered**

What does the user ultimately get told? What does a final report look like? How do you communicate effect size, uncertainty, inconclusive results, and the distinction between “no detectable benefit” and “evidence of no benefit”?

- At the end of an exploration the user sees whether it worked for them and how their results compare with others who did the same exploration. [BRIEF §2.4; WEB FAQ]
- We have developed an initial outline of the reports a user receives at different points in their exploration: a baseline summary, an interim report after each phase, a completion report, a personalised final report, and community comparisons (e.g. at weeks 2, 4, 6 and 8). We would like to develop these further with the Science Advisory Board. [REPORT pp.1–12]
- The example final report shows change against baseline (e.g. afternoon energy 5.1 → 6.6), each element’s effect with a strength label (“Moderate”, or “Experimental · too small to call”), a “keep list”, a comparison with the community, and a “Hold this lightly” message with limitations. [REPORT pp.6–8]
- The individual analysis is pre-set in each exploration’s analysis plan. Kind shares the statistical relationship between the choice and the outcome, using Bayesian and inferential analysis, and the user decides what to do with it. [BRIEF §10a; DECK s4, s14]

NOTE: How effect size, uncertainty and inconclusive results are presented, and how “no detectable benefit” is distinguished from “evidence of no benefit”, has not yet been defined. The example reports use qualitative strength labels but do not yet show uncertainty ranges.

SCIENCE ADVISORY BOARD INPUT REQUESTED: This is a key area where we are looking for support from the Science Advisory Board, as part of building the statistical analysis framework.

Q7 Safety monitoring and stopping rules **Partly answered**

What safety monitoring occurs during an experiment? What are the stopping rules, adverse-event procedures, escalation procedures, and responsibilities if a potentially dangerous symptom or finding appears?

- Adverse-event pathway: in-app detection (“report a concern” plus passive monitoring), then severity triage. Immediate risk leads to a non-dismissible redirect to 911 or 988; concerning but non-emergency issues go to a contracted, independent 24/7 clinical triage service (licensed professionals, Schmitt-Thompson protocols); anticipated effects are logged. [BRIEF §4.8; CAT liability tab, AE]
- By tier: Tier 2 adds a clinician-readable summary and aggregate pharmacovigilance-style review; Tier 3 adds a live triage handoff and a prompt to report to FDA MedWatch. [CAT liability tab; DECK s12]
- Responsibilities: the Safety Officer runs the pathway, sets and monitors pause triggers, and reports unanticipated problems to BRANY through the investigator of record (proposed within 5–10 business days). The investigator of record (the CSO) can pause research. [R&R pp.12–13; DECK s10]
- Safety signals raised in community features, which are moderated, are handed from the Product Team to the Safety Officer. [R&R p.9; WEB FAQ]

NOTE: Numeric pause thresholds for each exploration and stopping rules for individual users are not yet defined (the interim BRANY answers recommend proposing them). The Safety Officer role is not yet filled.

Q8 Regulatory framework **Answered**

What is the regulatory framework? When do you consider these activities clinical care versus human-subjects research? When would IRB review be required? Have you obtained legal/regulatory advice concerning FDA oversight, clinical decision-support rules, telehealth, and relevant state laws?

- Kind treats its activity (for launch) as general wellness and human-subjects research, not clinical care: it gives no diagnosis or treatment advice, and would “rather be reviewed than argue we need not be”. [BRIEF §4.1, §14; WEB FAQ]
- IRB review is required for every exploration before launch, through BRANY, under an umbrella protocol with templated sub-studies. The Common Rule is applied voluntarily; 21 CFR 50 and 56 will be considered for Tier 3. [BRIEF §13, §18–19; DECK s10–11; VB p2]
- FDA: the launch scope is intentionally outside device oversight under the General Wellness and Clinical Decision Support guidance (Jan 2026). Tier 3 follows FDA review, likely a 510(k); a 513(g) request is an option. [CAT liability tab, FDA; DECK s11]
- Legal advice: US counsel (Promise Legal, <https://promise.legal/>) is named for state licensure and liability. The FTC Health Breach Notification Rule and state health data laws (e.g. Washington MHMDA) are mapped with counsel. The angel round funds FDA, FTC, state and liability cover. [BRIEF §12, §16; VB p2; PITCH p15]

NOTE: We do not have, and are not expecting to get, any formal FDA position prior to launch, as the FDA issues ample guidance for self-assessment. No 513(g) request has been made and there is no 513(g) request planned for launch. Telehealth rules are not addressed as Kind does not provide any telehealth features. The interim BRANY answers are an analysis at this stage and will be built out and shared with legal counsel for their input.

Q9 Data use, ownership and consent **Partly answered**

What happens to the data? Who owns it, where is it stored, who can use it, and for what secondary purposes? Can it be pooled for research, publications, commercial partnerships, or AI development? What consent and privacy protections apply?

- Consent is layered: product consent; a separate, unbundled opt-in for research; Broad Consent for secondary analysis; specific consent for each exploration. Research consent can be withdrawn at any time; data already in completed analyses cannot be retrieved. [BRIEF §4.6, §17; WEB FAQ, privacy policy]
- Ownership and rights: the website says data is held in a personal data vault and “belongs to you”. Users can access, correct or delete their information and request a download of their exploration data. [WEB FAQ, privacy policy]
- De-identification uses the HIPAA Safe Harbor standard as a working definition, plus minimum cohort sizes. Researchers get aggregate outputs by default; row-level data only under a Data Use Agreement that bans re-identification and onward transfer. Kind does not sell data, has no ads, and never uses health data for advertising. [BRIEF §10d–e; CAT liability tab, FTC; WEB privacy policy]
- Pooling of de-identified data for research and preprint publication is intended. Researchers can access some de-identified data openly, more by subscription, and get priority access to studies they initiate. Other planned revenue (employer and insurer subsidies, brand partnerships) is described as operating under ethics, FDA, FTC and data-privacy cover. Change-of-control protections are proposed. [BRIEF §3, §10b–c; PITCH p11; VB p2; BRANY E]

NOTE: Assumption that the individual has access to, and can download, their own data, but that a version of that data can be used by Kind. Assumption that data will be stored in the US (note that the privacy policy currently says processing may take place in the US, UK and Netherlands; see inconsistency 5). The UK/EU GDPR workstream comes later.

NOTE: Rules on whether data may be used for AI or model development are yet to be agreed. The venture brief already describes “AI tools to aggregate data and identify patterns at scale” for researchers, so this needs deciding before launch.

Q10 Aggregating results across individuals **Partly answered**

How do you plan to aggregate results across individuals? Under what circumstances would n-of-1 trials be combined, and what is the intended inference—average treatment effect, heterogeneity of treatment response, prediction for future individuals, or something else?

- Two kinds of group analysis are planned: across individuals who completed the same exploration and met data quality thresholds (did it work overall, and for which sub-groups), and across explorations (estimating causal effects of behaviour changes from real-world data). [BRIEF §10b–c]
- Analyses use pre-specified analysis plans or explicit emulative analysis. Both protocols and study results are published as preprints with researcher attribution. [BRIEF §10b; DECK s3; PITCH p11; WEB FAQ]
- Intended inference: average effects and heterogeneity of response by sub-group. Bayesian analysis is part of the model, supported by Prof Chris Schmid's expertise. Community averages are already shown to users (e.g. "Kind explorers improved afternoon energy by about 1.1 points on average"). [BRIEF §10b; DECK s8, s14; REPORT p8]
- Prediction for future individuals is not used in the product: the matching engine does not use other people's outcomes. [CAT AI Matching Logic tab]

NOTE: The statistical model (for example a hierarchical Bayesian model across N-of-1 series), the rules for when trials are combined, and the minimum cohort thresholds are not specified. The website's description of matching conflicts with the last bullet (see inconsistency 2).